

Diffusion-Based Data Augmentation for Imbalanced Medical Image Classification

MSML612 Group Project, Final Report

Team (2 members): Jack Maiorino, Rithvik Kommareddy

Date: August 2026

Abstract

We asked whether diffusion-generated images can improve a downstream classifier on the rare classes of HAM10000, a severely imbalanced dermoscopic benchmark. We trained a class-conditional 37.1M-parameter pixel-space DDPM from scratch on a leak-free lesion-level split and generated 1,000 images per class. In this class-balanced run, a nearest-neighbor check in two perceptual spaces, calibrated on real-to-real distances with whole-lesion exclusion, flags 97% of generated dermatofibroma (df) images as near-copies of specific training images; the flags trace to 81 of the 84 df training images. Flag rates generally rise across training checkpoints as train-referenced KID falls. A size-matched held-out analysis is less dramatic than the unmatched comparison: for df, median LPIPS is 0.134 to an equal-sized training reference and 0.155 to validation, and 66.0% of synthetic images are closer to train versus 63.7% for the real-data baseline. Filtering does not improve any of the seven held-out KID point estimates, although subset variability is too wide to establish each change inferentially. Across three classifier seeds, classical augmentation raises mean rare-class F1 from 0.343 to 0.573; duplicated real, unfiltered synthetic, and filtered synthetic data score 0.289, 0.299, and 0.290. These descriptive results provide no evidence of a downstream gain from either synthetic arm and do not establish equivalence between filtered synthetic and duplicated real data.

1. Problem Statement and Research Question

Clinical image datasets are long-tailed, and classifiers underperform exactly where reliable detection matters most. Classical augmentation only perturbs existing pixels; generative augmentation promises new samples. Diffusion models are the state of the art in image generation and provide native class-conditional control, which motivated our original question: can diffusion-generated images meaningfully improve a rare-class classifier? The project's findings sharpened that question into the one this report answers:

When a diffusion model is trained from scratch on tens of images per class, are its "synthetic" rare-class samples new images at all, and does anything they add to a classifier survive controlling for trivial oversampling?

2. Data

HAM10000 (Tschandl et al., 2018) contains 10,015 dermoscopic images in 7 classes, from 66.9% melanocytic nevi (nv) down to 1.1% dermatofibroma (df). Many lesions are photographed more than once (4,501 of 10,015 images share a lesion with another image), so a naive image-level split leaks near-duplicates across the train/test boundary: averaged over 200 random image-level splits, 34.9% of test images share a lesion with a training image, and the rate is worst on the minority classes (mel 62.1%, df 51.1%). We therefore split at the lesion level, stratified by class, 70/15/15, seed 612, keeping all images; an assertion verifies that no lesion crosses splits. Full details are in the interim report; the resulting counts:

Class	Train	Val	Test	Total
nv	4,679	1,015	1,011	6,705
mel	776	163	174	1,113
bkl	765	165	169	1,099
bcc	354	77	83	514
akiec	227	55	45	327
vasc	96	23	23	142
df	84	15	16	115
All	6,981	1,513	1,521	10,015

df, with 84 training images from 51 distinct lesions, is the stress test the study turns on.

3. Generative Model

The generator is a class-conditional DDPM (Ho et al., 2020) built on the diffusers `UNet2DModel` (von Platen et al., 2022), trained from scratch in pixel space at 64x64: 37.1M parameters, channels 128/256/256/256 with self-attention at 16x16 and 8x8, a cosine noise schedule with $T = 1000$ (Nichol and Dhariwal, 2021), and classifier-free guidance (Ho and Salimans, 2021) with label dropout 0.1. One shared model serves all classes so the rare classes can reuse representations learned from the common ones, and training uses class-balanced sampling so they receive equal gradient exposure. As run: 100,000 steps of AdamW (Loshchilov and Hutter, 2019) at $1e-4$, effective batch 64 (batch 32 with 2-step gradient accumulation, which avoids VRAM spill on our 12 GB GPU), EMA decay 0.9999, 13.0 hours on a single GPU. Sampling uses DDIM (Song et al., 2021) at 50 steps with guidance 2.0 and seed 612; we generated 1,000 images per class.

DDPM samples at guidance 2.0




Figure 1. EMA samples at 100k steps, guidance 2.0. Legible lesion morphology in every class; the question the rest of the report answers is where that morphology comes from.

4. Generative Quality by the Standard Metrics

Per-class FID and KID against the real training images, with a finite-sample real-split reference formed by splitting each class's training set in half and scoring the halves. This reference is not a floor: it uses two half-sized real samples, whereas the synthetic comparison uses the full real class and 1,000 generated images, so the estimators have different finite-sample behavior. KID is reported $\times 1000$. For four classes (akiec, bcc, df, and vasc) the real-split FID reference exceeds the synthetic score, so FID is uninformative there and we rank those classes by KID only.

Class	FID	FID real-split ref.	KID	KID real-split ref.
akiec	46.52	70.96	14.10	-0.35
bcc	53.91	62.39	24.04	0.53
bkl	59.58	41.07	29.42	-0.15
df	38.33	107.42	5.68	3.75
mel	61.61	42.12	33.65	1.07
nv	52.30	11.14	24.56	-0.04
vasc	43.40	105.67	4.14	-5.27

The df and vasc reference entries also carry no variance: their halved train sets (42 and 48 images) sit below the 100-image KID subset cap, so every real-split subset is a permutation of the same images and the reported reference is a single draw, not a tight floor.

Taken at face value the two rarest classes, df and vasc, are the model's best. That reading is wrong, and the reason is the subject of Section 5: a generator that copies its training set can receive a low train-referenced KID by construction. These numbers are kept here because they are the metrics the field defaults to, and because the gap between this table and Section 5 is itself a finding: train-referenced fidelity metrics cannot be read as generative quality without a memorization check next to them.

5. Memorization

5.1 Detection method

For every generated image we compute the nearest-neighbor distance to its class's real training images in two spaces: LPIPS (Zhang et al., 2018), computed through an exact algebraic decomposition that reduces it to squared Euclidean distance on precomputed vectors, and cosine distance on unit-normalized InceptionV3 pool features. An image is flagged when it falls below the 5th percentile of that class's real-to-real leave-one-out nearest-neighbor distribution in either space. Flagged images are quarantined, not deleted.

Calibration matters more than the metric. HAM10000's repeated lesion views mean a plain leave-one-out null still contains sibling images of the same lesion, which drags the null distances down and makes the detector too lenient. We therefore recalibrated by excluding the entire lesion, not just the identical image, from each real image's neighbor set. This roughly doubles most thresholds (akiec

LPIPS 0.033 to 0.092). All headline numbers below use the lesion-calibrated thresholds; the union of two 5% rules has a nominal false-positive rate between 5 and 10%, which is the baseline to read the table against.

5.2 The model copies its rare classes

Class	Flagged (image null)	Flagged (lesion null)	Surviving images
df	93.4%	97.0%	30
vasc	89.8%	92.3%	77
akiec	50.1%	79.1%	209
bcc	35.4%	56.2%	438
bkl	4.1%	25.0%	750
mel	4.6%	16.9%	831
nv	0.7%	2.8%	972

Flagged fraction is strongly associated with inverse class size. This is consistent with repeated exposure: under class-balanced sampling over 100k steps, each df training image was presented roughly 11,000 times. Class size and per-image exposure are coupled in this design, however, so this run does not identify the cause of the association. Visual inspection of the worst and near-threshold examples supports the copy interpretation and suggests that the threshold can miss visually similar pairs (Figure 2). Only nv, with 4,679 training images, stays below the nominal false-positive band.




Figure 2. The lowest-distance flagged pairs: each column pairs a generated image with its nearest real training image. These are near-pixel copies.

The copying is broad, not collapse onto a few prototypes: df's flags trace back to 81 of its 84 training images (50 of 51 lesions), vasc's to 88 of 96, akiec's to 163 of 227, and the most-copied df source accounts for only 8% of flags. Thus the detected copying spans most of the df and vasc training sets rather than a few sources.

5.3 Copying rises as train-referenced KID falls

We scored checkpoints from 10k to 100k steps, plus a guidance 1.0 probe at 100k (Figure 3). Flag rates generally rise with training while train-referenced KID falls. But these two curves are not independent: copies are close to the empirical training distribution by definition, so train-KID improves as copying worsens. The trajectory is also not perfectly monotonic at the point level (the

70k checkpoint beats 90k on both axes, and guidance 1.0 cuts df's flag rate from 97% to 84% at similar KID), so sampling choices move the numbers. What no tested setting produced is an operating point that is simultaneously novel and faithful: early checkpoints are novel but visibly unconverged (KID 52 to 106), late ones are faithful by copying.




Figure 3. Left: lesion-calibrated flag rate. Right: KID x1000 vs train. Open markers: guidance 1.0 at 100k. Copying and train-referenced "quality" improve together, which is exactly why train-referenced quality cannot be trusted here.

5.4 Held-out comparison with size-matched references

A nearest-neighbor comparison is sensitive to the number of candidate references. The training sets are larger than validation, so we repeatedly subsample each training class to the validation class size before comparing distances. The table reports median LPIPS and the fraction closer to the matched training reference, averaged over 20 seeded subsamples. For the real-data baseline, each validation image plays the role of a generated image; same-lesion validation neighbors are excluded.

Class	Pool to train, matched	Pool to val	Pool closer to train	Real val to train, matched	Real val to val	Real closer to train
df	0.134	0.155	66.0%	0.149	0.160	63.7%
vasc	0.112	0.153	73.1%	0.155	0.173	45.0%
akiec	0.122	0.135	64.6%	0.141	0.145	42.8%
bcc	0.137	0.145	61.4%	0.141	0.143	52.2%
bkl	0.145	0.148	55.8%	0.145	0.148	51.6%
mel	0.151	0.160	62.8%	0.146	0.149	57.6%
nv	0.125	0.125	52.1%	0.096	0.096	48.5%

Size matching changes the interpretation. For df, the train-to-validation distance ratio is 1.15 rather than 20.2, and the closer-to-train fraction is only 2.3 percentage points above the real baseline. For vasc the ratio is 1.37 and the fraction is 28.1 points above baseline; akiec also shows a substantial fraction difference. These held-out results are descriptive and mixed, not an independent twenty-fold

copying result. The lesion-calibrated flag rates, source coverage, and inspected pairs in Section 5.2 remain the direct evidence of copying.

5.5 Filtering does not improve the held-out KID point estimates

The obvious salvage is to keep only the unflagged images. Measured against validation, all seven accepted-pool KID point estimates are higher than their unfiltered counterparts. The table reports KID x1000 mean +/- subset standard deviation, with real train scored against validation as a reference:

Class	Real train	All synthetic	Accepted only	Size-matched near-dup rate (all / accepted)
akiec	1.5 +/- 1.7	24.0 +/- 5.8	59.4 +/- 7.4	80% / 78%
bcc	4.1 +/- 2.7	43.6 +/- 7.8	69.9 +/- 7.5	72% / 68%
bkl	0.9 +/- 1.9	36.6 +/- 4.0	38.8 +/- 3.6	60% / 54%
df	6.7 +/- 9.0	15.0 +/- 14.6	26.1 +/- 10.3	76% / 30%
mel	1.3 +/- 1.8	40.5 +/- 5.4	43.0 +/- 5.2	49% / 44%
nv	-0.1 +/- 1.9	26.6 +/- 5.5	28.3 +/- 5.7	3% / 2%
vasc	8.7 +/- 6.5	21.1 +/- 7.6	33.7 +/- 9.0	73% / 62%

The held-out means support two observations. First, df's KID is 15.0 against validation rather than 5.7 against train, consistent with train-referenced KID rewarding copies. Second, filtering does not reveal a better-fidelity subset in these point estimates. The subset standard deviations are neither paired confidence intervals nor hypothesis tests; for example, df changes by 11.1 while the two subset standard deviations are 14.6 and 10.3. We therefore do not claim that filtering conclusively worsens every class.

The final column substantially reduces the original pool-size confound. For each of 20 repeats, the synthetic pool and real calibration set are subsampled to the same size: the smaller of the pool size, class training size, and 500. The LPIPS threshold is recomputed on that real subset with whole-lesion exclusion before the synthetic subset is scored. The matched real null rate is 4.8 to 5.3% except for accepted df, whose 30-image quantile yields 6.7%. The rare-class pool rates therefore remain far above their matched nulls after matching the number of images evaluated. They are detector rates under this LPIPS rule, not estimates of duplicate prevalence.

6. Downstream Comparison

The generative analysis motivates comparing synthetic data with a duplicated-real control. The classifier architecture follows the interim plan: a ResNet-18 (He et al., 2016) trained from scratch at 64x64 with class-balanced sampling. Architecture, AdamW optimizer, learning rate schedule, 6,000-step budget, evaluation cadence, and seed IDs are fixed across regimes. The arm manifest or train-time augmentation is the regime-defining difference. Five regimes were run for three seeds each (612, 613, 614):

1. Real data only.

2. Real + classical augmentation (flips, rotation, color jitter).
3. Real + duplicated-real oversampling (matching regime 5's added volume with copies of real images: the control the memorization result makes essential).
4. Real + unfiltered synthetic (1,000 per class).
5. Real + filtered synthetic (the accepted pools of Section 5.5).

The real-only and classical arms use the 6,981-row real manifest, duplicated-real and filtered-synthetic each add 3,307 class-matched rows, and unfiltered synthetic adds 7,000. These are fixed manifests, not validation-selected mixing ratios. Validation macro-F1 selects the checkpoint only; its test set is then evaluated once for that seed and regime. The primary summary is mean F1 on akiec, df, and vasc, with macro-F1 and balanced accuracy secondary. Results below are mean +/- sample standard deviation across three seeds.

Regime	Rare-class F1 (delta vs real)	Macro F1	Balanced acc.
1 Real only	0.343 +/- 0.016 (reference)	0.438 +/- 0.014	0.422 +/- 0.027
2 Classical aug	0.573 +/- 0.040 (+0.231)	0.604 +/- 0.025	0.611 +/- 0.018
3 Duplicated-real	0.289 +/- 0.011 (-0.054)	0.420 +/- 0.002	0.446 +/- 0.021
4 Unfiltered synthetic	0.299 +/- 0.022 (-0.044)	0.432 +/- 0.012	0.414 +/- 0.018
5 Filtered synthetic	0.290 +/- 0.048 (-0.052)	0.416 +/- 0.029	0.399 +/- 0.033

Regime	akiec F1	df F1	vasc F1
1 Real only	0.232 +/- 0.026	0.068 +/- 0.072	0.727 +/- 0.066
2 Classical aug	0.487 +/- 0.005	0.395 +/- 0.079	0.838 +/- 0.035
3 Duplicated-real	0.184 +/- 0.012	0.111 +/- 0.033	0.571 +/- 0.013
4 Unfiltered synthetic	0.181 +/- 0.034	0.060 +/- 0.053	0.656 +/- 0.073
5 Filtered synthetic	0.118 +/- 0.069	0.032 +/- 0.055	0.721 +/- 0.113

Classical augmentation has the highest three-seed mean for every rare class and raises mean rare-class F1 by 0.231 over real only. Neither synthetic arm improves the corresponding mean. Filtered synthetic and duplicated real have nearly equal aggregate rare-class means, a difference of +0.0016, but this is cancellation rather than demonstrated equivalence: filtered-minus-duplicated differences are -0.066 for akiec, -0.080 for df, and +0.150 for vasc. Their paired rare-class differences by seed are +0.0248, -0.0637, and +0.0436. With only three seeds and no prediction-level bootstrap or predeclared equivalence margin, these numbers show no observed synthetic-data gain in this experiment, not that filtered synthetic and duplicated real are interchangeable in the population. Per-run values are in `reports/classifier_runs.csv`, with arm summaries in `reports/classifier_summary.csv`.

7. Discussion

Plausible mechanism. Class-balanced sampling, our deliberate defense against rare-class neglect, multiplied each df image's expected exposure by roughly 12x relative to natural frequency, to about

11,000 presentations over the run. This makes repeated exposure a plausible contributor to the inverse class-size pattern. Exposure and class size are coupled, however, and there is no natural-frequency control, so this study cannot identify balancing as the cause. The supported scope is narrower: in this one 37M-parameter from-scratch DDPM run, the smallest classes have the highest calibrated detection rates.

Why we did not retrain. The candidate mitigations, in the order we would try them, are capping the oversampling weight (e.g. 90% empirical + 10% balanced sampling, which cuts per-image exposure by 5 to 10x), dihedral train-time augmentation with an augmentation-aware detector, a smaller model, scheduled guidance, and memorization-based early stopping. Each costs a full 13-hour training run plus re-scoring, and a single unreplicated attempt would support weaker claims than the analysis above. We chose to spend the remaining budget characterizing the negative result and measuring its downstream cost. The honest statement of scope: across the tested checkpoints, guidance scales, and sampler, no acceptable operating point was observed; the experiment does not estimate true copy prevalence (the detector's null rate is 5 to 10%, and it under-flags visually confirmed near-copies), identify a single cause, or rule out untested mitigations.

Limitations. Everything is established at 64x64 on one generator run. The validation references for the rarest classes are small (df 15, vasc 23 images), so held-out KID has substantial subset variability. Size matching also shows that the earlier unmatched train-versus-validation distance contrast was mostly a reference-count effect for df. Classifier results cover three seeds and have no inferential or equivalence analysis. Flag rates are detection rates under a calibrated threshold, not prevalence estimates. The ethical stake is worth stating plainly: in a medical setting, lightly perturbed copies of real patient images can create privacy risk while being presented as augmentation, so a memorization check should accompany such a pipeline. The run JSONs did not record immutable code and input hashes at execution time; the supplied run manifest is a post-hoc artifact inventory, not execution attestation.

8. Deviations from the Interim Plan

- **The StyleGAN2-ADA baseline was dropped.** Once the memorization result landed, the informative comparison became synthetic-vs-duplicated-real, not diffusion-vs-GAN. A GAN trained on the same data would require its own memorization and downstream evaluation, and the available compute went to the checkpoint sweep and held-out analysis instead. No GAN evidence was collected, so this report makes no diffusion-vs-GAN claim.
- **Four regimes became five.** The replication control the interim report mentioned in passing (duplicated-real oversampling) was promoted to the decisive comparator, and the diffusion regime split into unfiltered and filtered variants.
- **Classifier replication and inference were reduced.** The interim protocol specified five paired seeds, per-class Wilcoxon tests with Holm correction, and prediction-level bootstrap intervals. Execution used three seeds and did not save prediction-level outputs, so none of those inferential analyses was performed. The report therefore treats the classifier results as descriptive means and standard deviations. Moreover, a two-sided exact signed-rank test with only five

nonzero pairs could not attain $p < 0.05$, so the originally specified seed count was itself insufficient for that test.

- **Ratio selection and planned ablations were omitted.** The executed arms use one fixed manifest each; validation selects checkpoints, not mixing ratios. The 0/25/50/100% ratio sweep and the 25/50/100% generator sample-size study were not run.
- **The guidance sweep narrowed.** Instead of {1.0, 2.0, 3.0} for sample quality, we probed guidance 1.0 specifically for its effect on copying (Figure 3).

9. Related Work

The diffusion foundations are as in the interim report: DDPM (Ho et al., 2020), the cosine schedule (Nichol and Dhariwal, 2021), classifier-free guidance (Ho and Salimans, 2021), and the diffusers implementation (von Platen et al., 2022). Frid-Adar et al. (2018) and Trabucco et al. (2024) report generative-augmentation gains; our result is a caution about the regime where those gains are claimed with the least data. On measurement, FID is due to Heusel et al. (2017), KID to Binkowski et al. (2018), and LPIPS to Zhang et al. (2018). Our held-out comparison follows the train-versus-test distance intuition of Meehan et al. (2020), but it does not implement their formal three-sample statistic or hypothesis test. Carlini et al. (2023) extracted training images from large diffusion models; Dar et al. (2025) documented the same phenomenon in medical diffusion models, including copies that survive augmented training, which informed our choice not to treat augmentation alone as a guaranteed fix.

10. References

1. Ho, J., Jain, A., and Abbeel, P. (2020). *Denoising Diffusion Probabilistic Models*. NeurIPS.
2. Nichol, A., and Dhariwal, P. (2021). *Improved Denoising Diffusion Probabilistic Models*. ICML.
3. Ho, J., and Salimans, T. (2021). *Classifier-Free Diffusion Guidance*. NeurIPS Workshop.
4. Tschandl, P., Rosendahl, C., and Kittler, H. (2018). *The HAM10000 Dataset*. Scientific Data.
5. Frid-Adar, M., et al. (2018). *GAN-based Synthetic Medical Image Augmentation for Increased CNN Performance in Liver Lesion Classification*. Neurocomputing.
6. Trabucco, B., et al. (2024). *Effective Data Augmentation with Diffusion Models*. ICLR.
7. Heusel, M., et al. (2017). *GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium*. NeurIPS.
8. Binkowski, M., et al. (2018). *Demystifying MMD GANs*. ICLR.
9. Zhang, R., et al. (2018). *The Unreasonable Effectiveness of Deep Features as a Perceptual Metric*. CVPR.
10. Meehan, C., Chaudhuri, K., and Dasgupta, S. (2020). *A Non-Parametric Test to Detect Data-Copying in Generative Models*. AISTATS.
11. Carlini, N., et al. (2023). *Extracting Training Data from Diffusion Models*. USENIX Security.
12. Dar, S. U. H., et al. (2025). *Unconditional Latent Diffusion Models Memorize Patient Imaging Data*. Nature Biomedical Engineering.
13. Karras, T., et al. (2020). *Training Generative Adversarial Networks with Limited Data (StyleGAN2-ADA)*. NeurIPS.
14. von Platen, P., et al. (2022). *Diffusers: State-of-the-Art Diffusion Models*. GitHub.
<https://github.com/huggingface/diffusers>
15. He, K., Zhang, X., Ren, S., and Sun, J. (2016). *Deep Residual Learning for Image Recognition*. CVPR.
16. Loshchilov, I., and Hutter, F. (2019). *Decoupled Weight Decay Regularization*. ICLR.
17. Song, J., Meng, C., and Ermon, S. (2021). *Denoising Diffusion Implicit Models*. ICLR.